

In [1]:

```
1 import pandas as pd
```

In [2]:

```
1 df = pd.read_csv('LoanApprovalPrediction.csv')
```

In [3]:

```
1 df.head()
```

Out[3]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credi
0	LP001002	Male	No	0.0	Graduate	No	5849	0.0	NaN	360.0	
1	LP001003	Male	Yes	1.0	Graduate	No	4583	1508.0	128.0	360.0	
2	LP001005	Male	Yes	0.0	Graduate	Yes	3000	0.0	66.0	360.0	
3	LP001006	Male	Yes	0.0	Not Graduate	No	2583	2358.0	120.0	360.0	
4	LP001008	Male	No	0.0	Graduate	No	6000	0.0	141.0	360.0	

In [4]:

```
1 df.shape
```

Out[4]:

(598, 13)

In [5]:

```
1 df.info
```

Out[5]:

```
<bound method DataFrame.info of
0  LP001002  Male  No  0.0  Graduate  No
1  LP001003  Male  Yes  1.0  Graduate  No
2  LP001005  Male  Yes  0.0  Graduate  Yes
3  LP001006  Male  Yes  0.0  Not Graduate  No
4  LP001008  Male  No  0.0  Graduate  No
..  ...  ...  ...  ...  ...  ...
593 LP002978  Female  No  0.0  Graduate  No
594 LP002979  Male  Yes  3.0  Graduate  No
595 LP002983  Male  Yes  1.0  Graduate  No
596 LP002984  Male  Yes  2.0  Graduate  No
597 LP002990  Female  No  0.0  Graduate  Yes

ApplicantIncome  CoapplicantIncome  LoanAmount  Loan_Amount_Term  \
0      5849      0.0      NaN      360.0
1      4583     1508.0     128.0     360.0
2      3000      0.0      66.0     360.0
3      2583     2358.0     120.0     360.0
4      6000      0.0     141.0     360.0
..      ...      ...      ...      ...
593     2900      0.0      71.0     360.0
594     4106      0.0      40.0     180.0
595     8072     240.0     253.0     360.0
596     7583      0.0     187.0     360.0
597     4583      0.0     133.0     360.0

Credit_History  Property_Area  Loan_Status
0      1.0      Urban      Y
1      1.0      Rural      N
2      1.0      Urban      Y
3      1.0      Urban      Y
4      1.0      Urban      Y
..      ...      ...      ...
593     1.0      Rural      Y
594     1.0      Rural      Y
595     1.0      Urban      Y
596     1.0      Urban      Y
597     0.0  Semiurban      N
```

[598 rows x 13 columns]>

In [6]:

```
1 df.describe()
```

Out[6]:

	Dependents	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
count	586.000000	598.000000	598.000000	577.000000	584.000000	549.000000
mean	0.755973	5292.252508	1631.499866	144.968804	341.917808	0.843352
std	1.007751	5807.265364	2953.315785	82.704182	65.205994	0.363800
min	0.000000	150.000000	0.000000	9.000000	12.000000	0.000000
25%	0.000000	2877.500000	0.000000	100.000000	360.000000	1.000000
50%	0.000000	3806.000000	1211.500000	127.000000	360.000000	1.000000
75%	1.750000	5746.000000	2324.000000	167.000000	360.000000	1.000000
max	3.000000	81000.000000	41667.000000	650.000000	480.000000	1.000000

In [7]:

```
1 df.isna().sum()
```

Out[7]:

Loan_ID	0
Gender	0
Married	0
Dependents	12
Education	0
Self_Employed	0
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	21
Loan_Amount_Term	14
Credit_History	49
Property_Area	0
Loan_Status	0
dtype:	int64

In [8]:

```
1 #check the uniqueness of Loan Id column
2 df.Loan_ID.nunique()
```

Out[8]:

598

In [9]:

```
1 #drop the Loan_ID column as it is not required (no duplicates)
2 df.drop(['Loan_ID'], axis = 1, inplace = True)
```

In [10]:

```
1 df.head()
```

Out[10]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	F
0	Male	No	0.0	Graduate	No	5849	0.0	NaN	360.0	1.0	
1	Male	Yes	1.0	Graduate	No	4583	1508.0	128.0	360.0	1.0	
2	Male	Yes	0.0	Graduate	Yes	3000	0.0	66.0	360.0	1.0	
3	Male	Yes	0.0	Not Graduate	No	2583	2358.0	120.0	360.0	1.0	
4	Male	No	0.0	Graduate	No	6000	0.0	141.0	360.0	1.0	

In [11]:

```
1 #Total missing value cells
2 df.isna().sum().sum()
```

Out[11]:

96

In [12]:

```
1 df.dtypes
```

Out[12]:

Gender object
Married object
Dependents float64
Education object
Self_Employed object
ApplicantIncome int64
CoapplicantIncome float64
LoanAmount float64
Loan_Amount_Term float64
Credit_History float64
Property_Area object
Loan_Status object
dtype: object

In [13]:

```
1 #convert Gender column from object to int datatype  
2 df.Gender = df.Gender.map({'Male' : 0, 'Female' : 1})
```

In [14]:

```
1 df.head()
```

Out[14]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	F
0	0	No	0.0	Graduate	No	5849	0.0	NaN	360.0	1.0	
1	0	Yes	1.0	Graduate	No	4583	1508.0	128.0	360.0	1.0	
2	0	Yes	0.0	Graduate	Yes	3000	0.0	66.0	360.0	1.0	
3	0	Yes	0.0	Not Graduate	No	2583	2358.0	120.0	360.0	1.0	
4	0	No	0.0	Graduate	No	6000	0.0	141.0	360.0	1.0	

In [20]:

```
1 #Convert all categorical(object) columns using LabelEncoder  
2 from sklearn.preprocessing import LabelEncoder  
3 label_encoder = LabelEncoder()  
4 obj = (df.dtypes == 'object')  
5 print(list(obj[obj].index))#List of categorical objects  
6 for col in list(obj[obj].index):  
7     df[col] = label_encoder.fit_transform(df[col])
```

['Married', 'Education', 'Self_Employed', 'Property_Area', 'Loan_Status']

In [21]:

```
1 df.head()
```

Out[21]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	F
0	0	0	0.0	0	0	5849	0.0	NaN	360.0	1.0	
1	0	1	1.0	0	0	4583	1508.0	128.0	360.0	1.0	
2	0	1	0.0	0	1	3000	0.0	66.0	360.0	1.0	
3	0	1	0.0	1	0	2583	2358.0	120.0	360.0	1.0	
4	0	0	0.0	0	0	6000	0.0	141.0	360.0	1.0	

In [22]:

```
1 df.dtypes
```

Out[22]:

```
Gender          int64
Married         int32
Dependents      float64
Education       int32
Self_Employed   int32
ApplicantIncome int64
CoapplicantIncome float64
LoanAmount      float64
Loan_Amount_Term float64
Credit_History  float64
Property_Area   int32
Loan_Status     int32
dtype: object
```

In [23]:

```
1 df.columns
```

Out[23]:

```
Index(['Gender', 'Married', 'Dependents', 'Education', 'Self_Employed',
      'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
      'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'],
      dtype='object')
```

In [24]:

```
1 #fill in the missing rows with mean
2 for i in df.columns:
3     df[i] = df[i].fillna(df[i].mean())
```

In [25]:

```
1 df.isna().sum()
```

Out[25]:

```
Gender          0
Married         0
Dependents      0
Education       0
Self_Employed   0
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount      0
Loan_Amount_Term 0
Credit_History  0
Property_Area   0
Loan_Status     0
dtype: int64
```

Training the model

In [26]:

```
1
2 X= df.drop(['Loan_Status'],axis=1)
3 y = df['Loan_Status']
```

In [27]:

```
1 X
```

Out[27]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
0	0	0	0.0	0	0	5849	0.0	144.968804	360.0	1.0
1	0	1	1.0	0	0	4583	1508.0	128.000000	360.0	1.0
2	0	1	0.0	0	1	3000	0.0	66.000000	360.0	1.0
3	0	1	0.0	1	0	2583	2358.0	120.000000	360.0	1.0
4	0	0	0.0	0	0	6000	0.0	141.000000	360.0	1.0
...
593	1	0	0.0	0	0	2900	0.0	71.000000	360.0	1.0
594	0	1	3.0	0	0	4106	0.0	40.000000	180.0	1.0
595	0	1	1.0	0	0	8072	240.0	253.000000	360.0	1.0
596	0	1	2.0	0	0	7583	0.0	187.000000	360.0	1.0
597	1	0	0.0	0	1	4583	0.0	133.000000	360.0	0.0

598 rows × 11 columns



In [28]:

```
1 y
```

Out[28]:

0	1
1	0
2	1
3	1
4	1
...	..
593	1
594	1
595	1
596	1
597	0

Name: Loan_Status, Length: 598, dtype: int32

In [48]:

```
1 #Split the data into train and test
2 from sklearn.model_selection import train_test_split
3 X_train,X_test,y_train,y_test = train_test_split(X, y ,test_size=0.3, random_state=7)
```

Model Selection

In [49]:

```
1 from sklearn.linear_model import LogisticRegression
2 from sklearn.linear_model import RidgeClassifier
3 from sklearn.tree import DecisionTreeClassifier
4 from sklearn.neighbors import KNeighborsClassifier
5 from sklearn.ensemble import RandomForestClassifier
6 from sklearn.metrics import accuracy_score
```

In [50]:

```
1 models = []
2 models.append(('Logistic Regression', LogisticRegression()))
3 models.append(('Ridge Classifier', RidgeClassifier()))
4 models.append(('Decision Tree Classifier', DecisionTreeClassifier()))
5 models.append(('K-Neighbors Classifier', KNeighborsClassifier()))
6 models.append(('Random Forest Classifier', RandomForestClassifier()))
```

In [51]:

```
1 def model_selection(model):
2     model.fit(X_train,y_train)
3     y_pred = model.predict(X_test)
4     return accuracy_score(y_test,y_pred)*100
```

In [52]:

```
1 for name,model in models:  
2     #print(name,model)  
3     print(f'{name} : {model_selection(model)}')
```

Logistic Regression : 81.66666666666667
Ridge Classifier : 82.22222222222221
Decision Tree Classifier : 75.55555555555556
K-Neighbors Classifier : 66.11111111111111
Random Forest Classifier : 80.55555555555556

Ridge Classifier Performs better

In []:

1